

Spam Email Detector Machine Learning Model

This is my project that I made during my free time to demonstrate knowledge in machine learning and cybersecurity. This model has been created to identify spam/phishing emails.

```
1 import pandas as pd
2 import numpy as np
3 from sklearn.model_selection import train_test_split
4 from sklearn.feature_extraction.text import TfidfVectorizer
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
7 import nltk
8 from nltk.corpus import stopwords
9 import re
10 from sklearn.model_selection import train_test_split
```

```
1 nltk.download('stopwords')

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
True
```

The file that contains around 5000 emails classified into spam or not spam.

```
1 data = pd.read_csv('https://github.com/yushika-j/Playing-with-datasets/raw/master/Email%20Spam%20Filtering/emails.csv', encoding='utf-8')
2 data.head()
```

	text	spam
0	Subject: naturally irresistible your corporate...	1
1	Subject: the stock trading gunslinger fanny i...	1
2	Subject: unbelievable new homes made easy im ...	1
3	Subject: 4 color printing special request add...	1
4	Subject: do not have money , get software cds ...	1

Next steps:

Generate code with data

View recommended plots

New interactive sheet

Data Preprocessing

Once loaded, I proceeded to preprocess and vectorize the text data to convert it into a format that a machine learning model can understand. I used TF-IDF or CountVectorizer from scikit-learn to convert the email content into a term-document matrix.

```
1 print(data.isnull().sum())
2
3 # Drop any rows with missing values (if any)
4 data.dropna(inplace=True)
5
6 # Check for duplicates
7 data.drop_duplicates(inplace=True)
```

```
[0] text      0
      spam    0
      dtype: int64
```

```
1 data['text'] = data['text'].str.lower()
2 #convert to lowercase
```

```
1 data['text'] = data['text'].apply(lambda x: re.sub(r'^a-z\s', '', x))
2 # remove special characters
```

```
1 #remove stopwords like 'the' and 'and'
2 stop_words = set(stopwords.words('english'))
3 data['text'] = data['text'].apply(lambda x: ' '.join([word for word in x.split() if word not in (stop_words)]))
```

Text Vectorization

To convert the cleaned text into a format that can be used by a machine learning model, use TF-IDF or Count Vectorizer to create a term-document matrix.

This will transform each email into a vector of numerical features where each feature will represent a frequency.

```
1 # Initialize the TF-IDF Vectorizer
2 vectorizer = TfidfVectorizer(max_features=1000) # Limit to 1000 most frequent words
3 X = vectorizer.fit_transform(data['text']).toarray() # Transforms text data to numerical array
4 # Target variable
5 y = data['spam']
6 #print the data
7 print(X)
```

```
[0. 0. 0. ... 0. 0. 0. ]
[0. 0. 0. ... 0. 0. 0. ]
[0. 0. 0. ... 0. 0. 0. ]
...
[0. 0. 0. ... 0. 0. 0. ]
[0. 0. 0. ... 0. 0. 0. ]
[0.15747477 0. 0.117228 ... 0. 0. 0. ]]
```

Train-Test split

```
1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
2
```

Train a Classification Model

I started with a model like Logistic Regression for simplicity purposes. If the requirements increase, I might change the model to a neural network model.

```
1 from sklearn.linear_model import LogisticRegression
2
3 model = LogisticRegression()
4 model.fit(X_train, y_train)
```

LogisticRegression

LogisticRegression()

Model Evaluation

```
1 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
2
3 y_pred = model.predict(X_test)
4
5 accuracy = accuracy_score(y_test, y_pred)
6 precision = precision_score(y_test, y_pred)
7 recall = recall_score(y_test, y_pred)
8 f1 = f1_score(y_test, y_pred)
9
10 print(f'Accuracy: {accuracy}')
11 print(f'Precision: {precision}')
12 print(f'Recall: {recall}')
13 print(f'F1 Score: {f1}')
14
```

Accuracy: 0.9754170324846356
Precision: 0.9855072463768116
Recall: 0.918918918918919
F1 Score: 0.951048951048951

Hyperparameter tuning

For optimal performance, I used parameter tuning to ensure the model works perfectly. In this case, I used GridSearchCV.

```
1 from sklearn.model_selection import GridSearchCV
2
3 param_grid = {'C': [0.1, 1, 10, 100]}
4 grid = GridSearchCV(LogisticRegression(), param_grid, refit=True, verbose=2)
5 grid.fit(X_train, y_train)
6
```

Fitting 5 folds for each of 4 candidates, totalling 20 fits

[CV] ENDC=0.1; total time= 0.1s

[CV] ENDC=0.1; total time= 0.1s

[CV] ENDC=0.1; total time= 0.1s

[CV] ENDC=0.1; total time= 0.1s

[CV] ENDC=0.1; total time= 0.1s

[CV] ENDC=1; total time= 0.1s

[CV] ENDC=1; total time= 0.1s

[CV] ENDC=1; total time= 0.1s

[CV] ENDC=1; total time= 0.1s

[CV] ENDC=1; total time= 0.1s

[CV] ENDC=10; total time= 0.1s

[CV] ENDC=10; total time= 0.1s

[CV] ENDC=10; total time= 0.1s

[CV] ENDC=10; total time= 0.2s

[CV] ENDC=100; total time= 0.1s

[CV] ENDC=100; total time= 0.1s

[CV] ENDC=100; total time= 0.2s

[CV] ENDC=100; total time= 0.1s

[CV] ENDC=100; total time= 0.2s

GridSearchCV

best_estimator_: LogisticRegression

LogisticRegression

The model applies a lower penalty for large coefficients, allowing it to fit the data more closely.

```
1 best_model = grid.best_estimator_
2 y_pred = best_model.predict(X_test)
```

Retrain on the full dataset

```
1 model = LogisticRegression(C=100)
2 model.fit(X_train, y_train)
```

LogisticRegression

LogisticRegression(C=100)

Cross Validation

Cross-validation is a technique used in machine learning to assess how well a model generalizes to unseen data. Instead of training and testing a model once on a single split of the data, cross-validation divides the data into multiple parts (called "folds") and trains and tests the model multiple times. This process provides a more reliable estimate of model performance and helps prevent overfitting.

```
1 # Convert cv_results_ to a DataFrame for easier viewing
2 cv_results = pd.DataFrame(grid.cv_results_)
3
4 # Display relevant columns
5 print("\nCross-Validation Results:")
6 print(cv_results[['param_C', 'mean_test_score', 'std_test_score']])
7
```



Cross-Validation Results:			
	param_C	mean_test_score	std_test_score
0	0.1	0.911764	0.010151
1	1.0	0.975199	0.006305
2	10.0	0.984637	0.003249
3	100.0	0.985076	0.003768